

# CPSC 481

# Artificial Intelligence

Mr. Ayush Bhardwaj

[abhardwaj@fullerton.edu](mailto:abhardwaj@fullerton.edu)

# What we will cover this week

- Adversarial search
  - Game trees
  - Zero-sum games
  - Minimax algorithm
  - Evaluation functions
  - $\alpha$ - $\beta$  pruning

# Game Playing

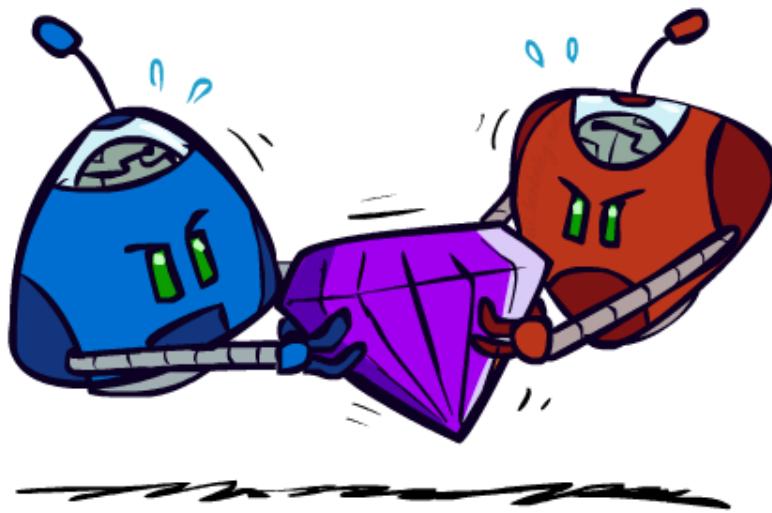
Why do AI researchers study game playing?

1. It's a good reasoning problem, formal and nontrivial.
2. Direct comparison with humans and other computer programs is easy.

# Behavior from Computation



# Adversarial Games

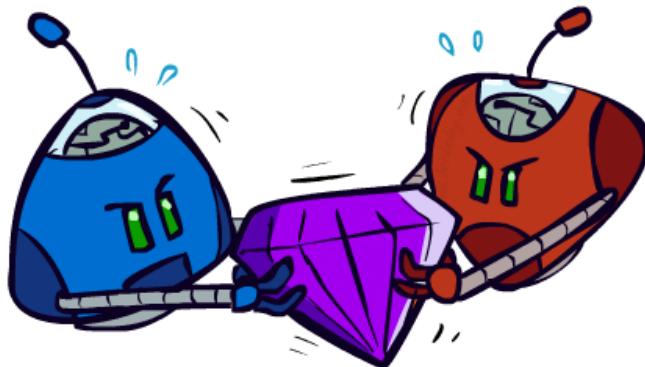


# Types of Games

- Many different kinds of games
  - Deterministic or stochastic?
  - One, two, or more players?
  - Zero sum?
  - Perfect information (can you see the state)?
- Want algorithms for calculating a **strategy (policy)** which recommends a move from each state



# Zero-Sum Games



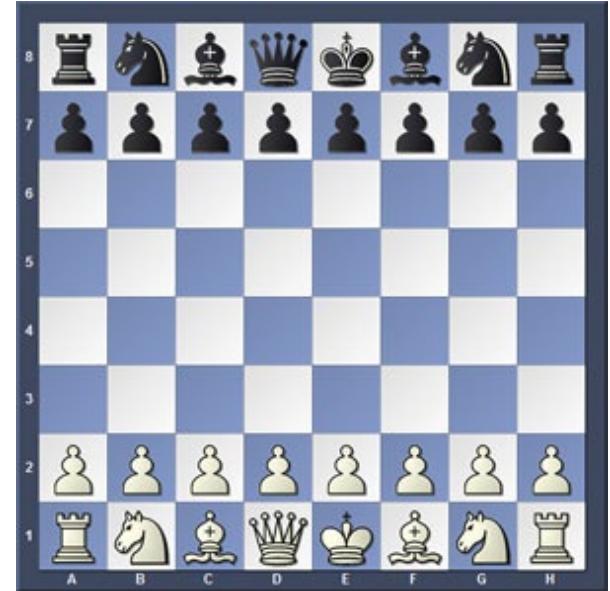
- Zero-Sum Games
  - Agents have opposite utilities (values on outcomes)
  - One player tries to maximize value, the other player minimizes it
  - a single value that one maximizes and the other minimizes
  - Adversarial, pure competition
- General Games
  - Agents have independent utilities
  - Cooperation, indifference, competition, and more are all possible

In this class, we will focus on:

- Deterministic, 2-player, zero-sum games with perfect information
- Examples: Tic-tac-toe, Chess, checkers, Othello

# Example: chess

- Players: Black, White
- State  $s$ : position of all pieces, whose turn it is
- Actions( $s$ ): legal chess moves that Player(s) can make
- IsTerminal( $s$ ): whether  $s$  is checkmate or draw
- Utility( $s$ ): +1 if white wins, 0 if draw, -1 if black wins

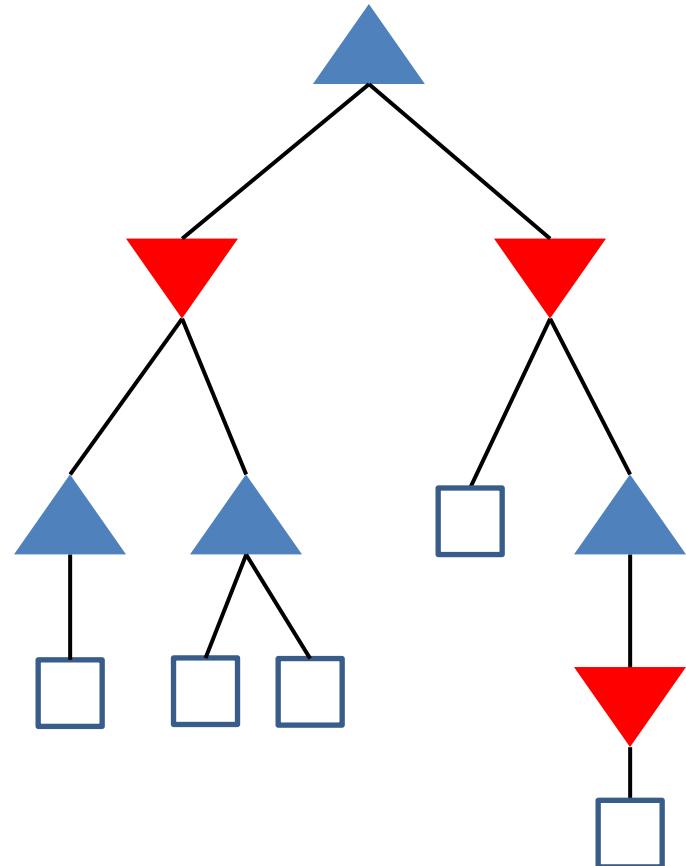


# Halving game

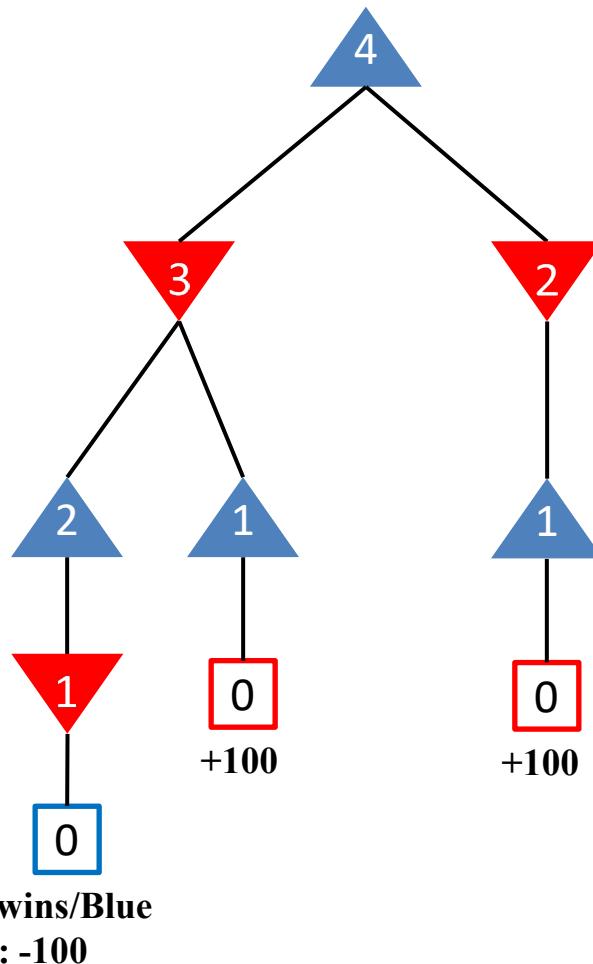
- Start with some number  $N$
- Players take turns. Two actions:
  - Decrement number, or
  - Replace with  $\text{floor}(N/2)$  (e.g.,  $\text{floor}(5/2) = 2$ )
- The player that makes it 0 wins
- A win is worth +100 points, loss is -100
- Start with  $N=4$
- What is the optimal action at the beginning?

# Game tree

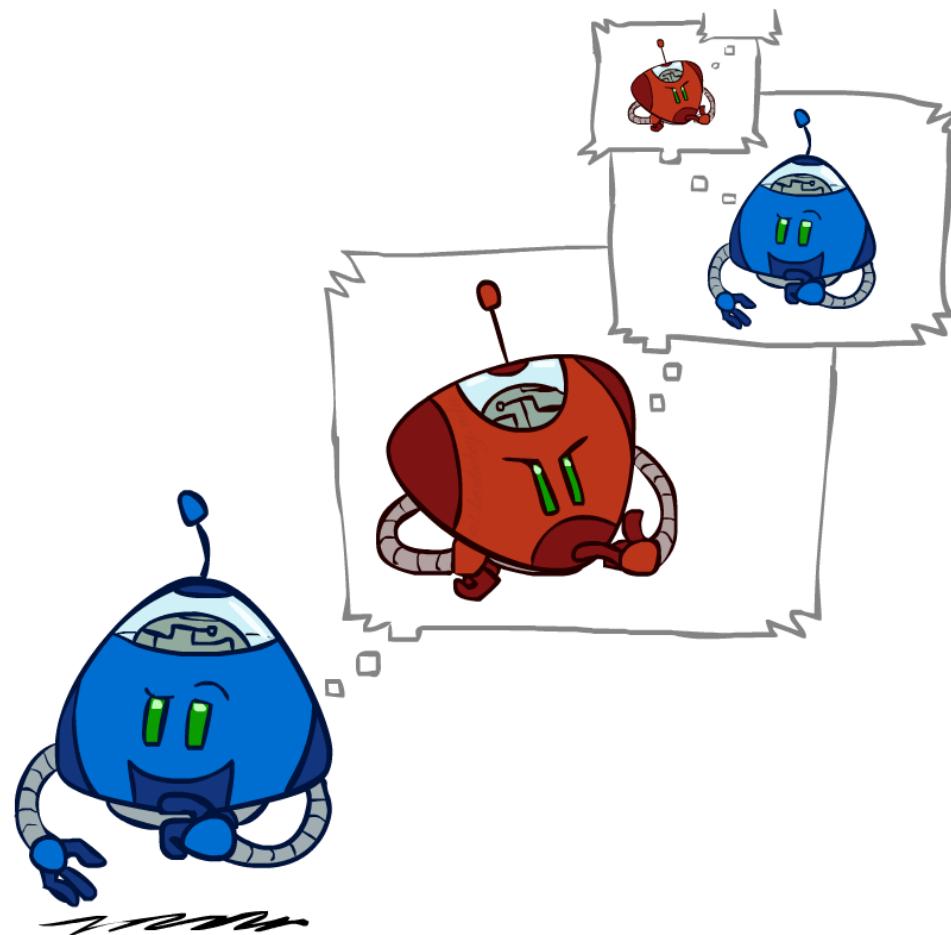
- Each node is a decision point for a player
- Decisions alternate at every level
  - Root: me (maximizing player)
  - 2<sup>nd</sup> level: opponent (minimizing player)
  - 3<sup>rd</sup> level: opponent (maximizing player again)
  - ...
- Leaves are terminal states: **outcome is known**
  - a win/loss/tie/points
  - Can assign a **utility value** from the point of view of the maximizing player
- Each root-to-leaf path is a possible outcome of the game



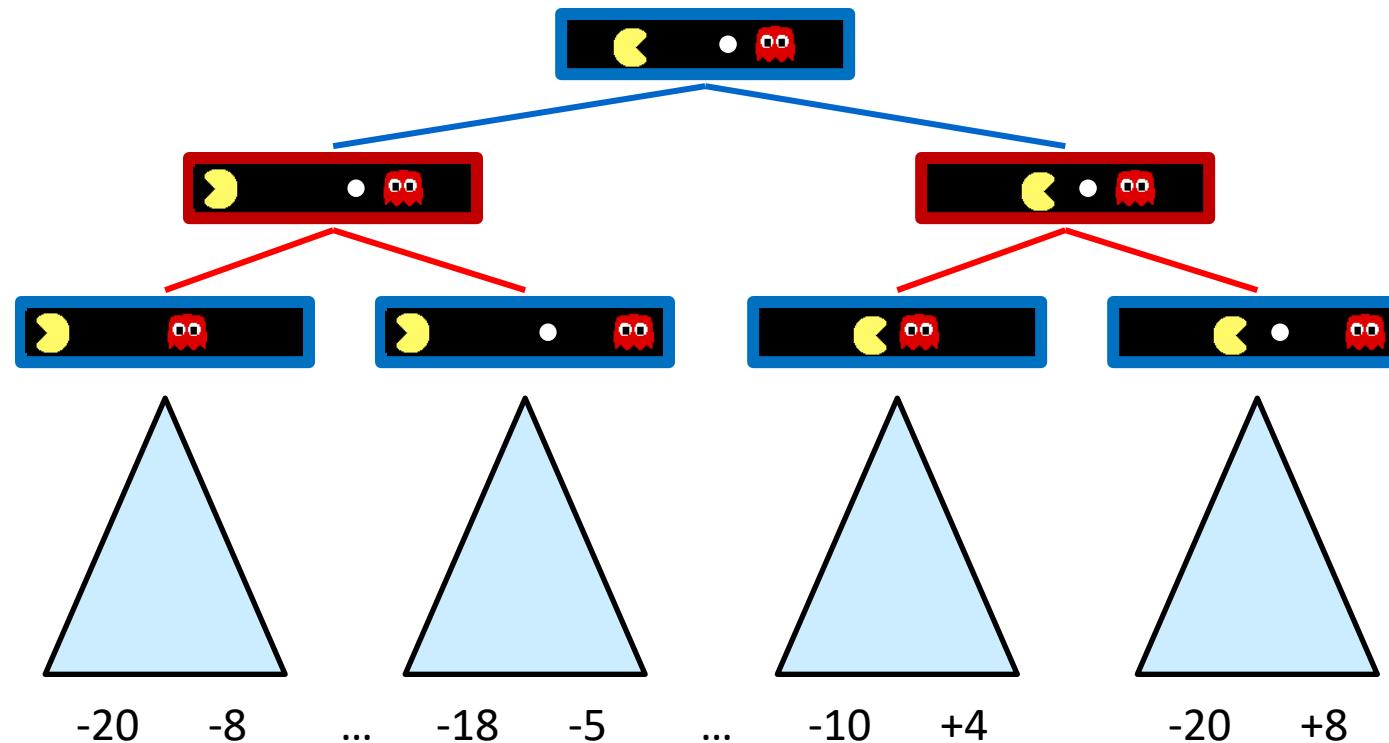
# Game tree for the halving game with N=4



# Adversarial Search



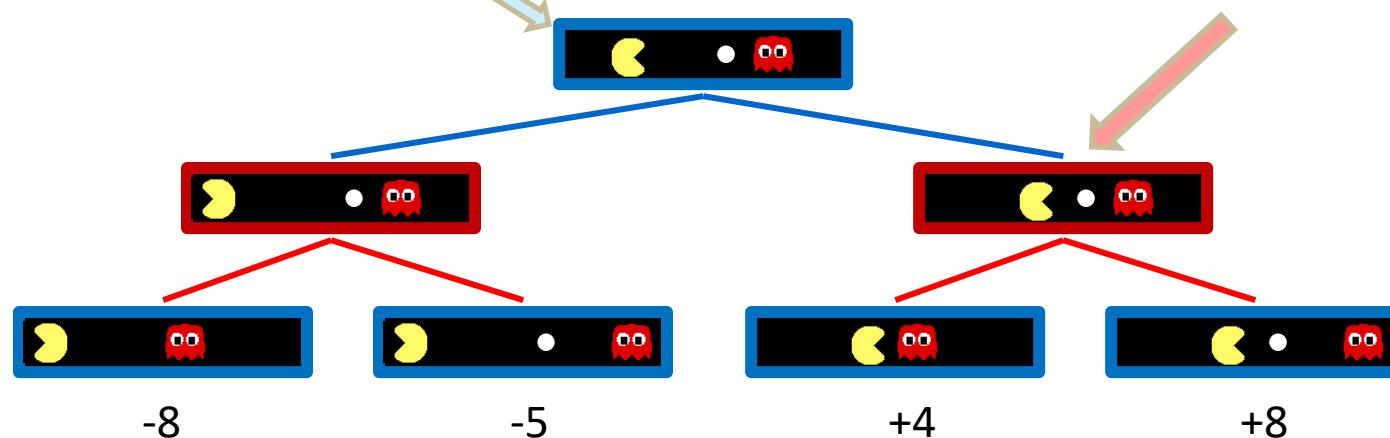
# Adversarial Game Trees



# Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$



States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

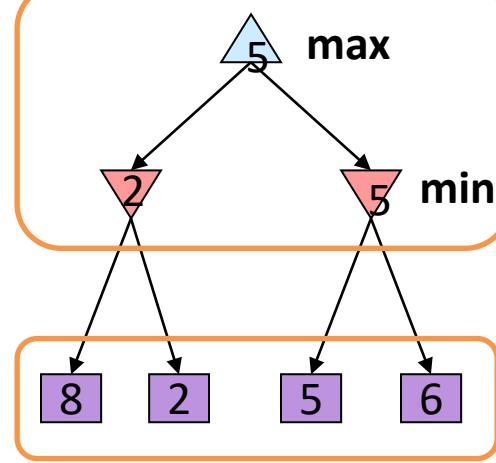
Terminal States:

$$V(s) = \text{known}$$

# Adversarial Search (Minimax)

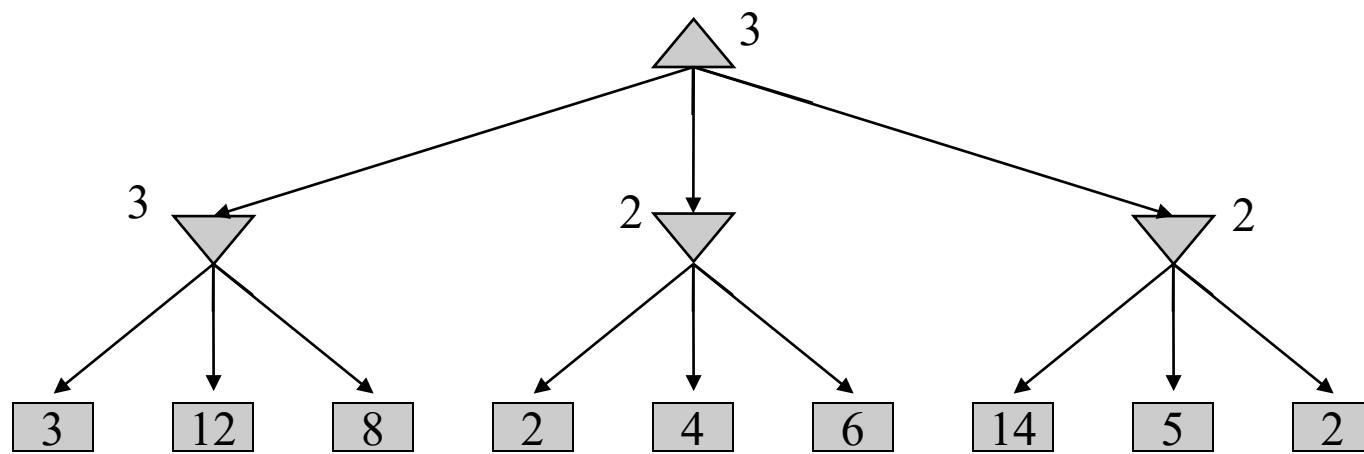
- Deterministic, zero-sum games:
  - Tic-tac-toe, chess, checkers
  - One player maximizes result
  - The other minimizes result
- Minimax search:
  - A state-space search tree
  - Players alternate turns
  - Compute each node's **minimax value**:  
the best achievable utility against a rational (optimal) adversary

**Minimax values:**  
computed recursively



**Terminal values:**  
part of the game

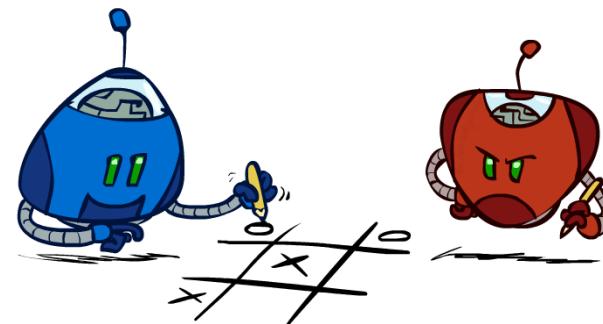
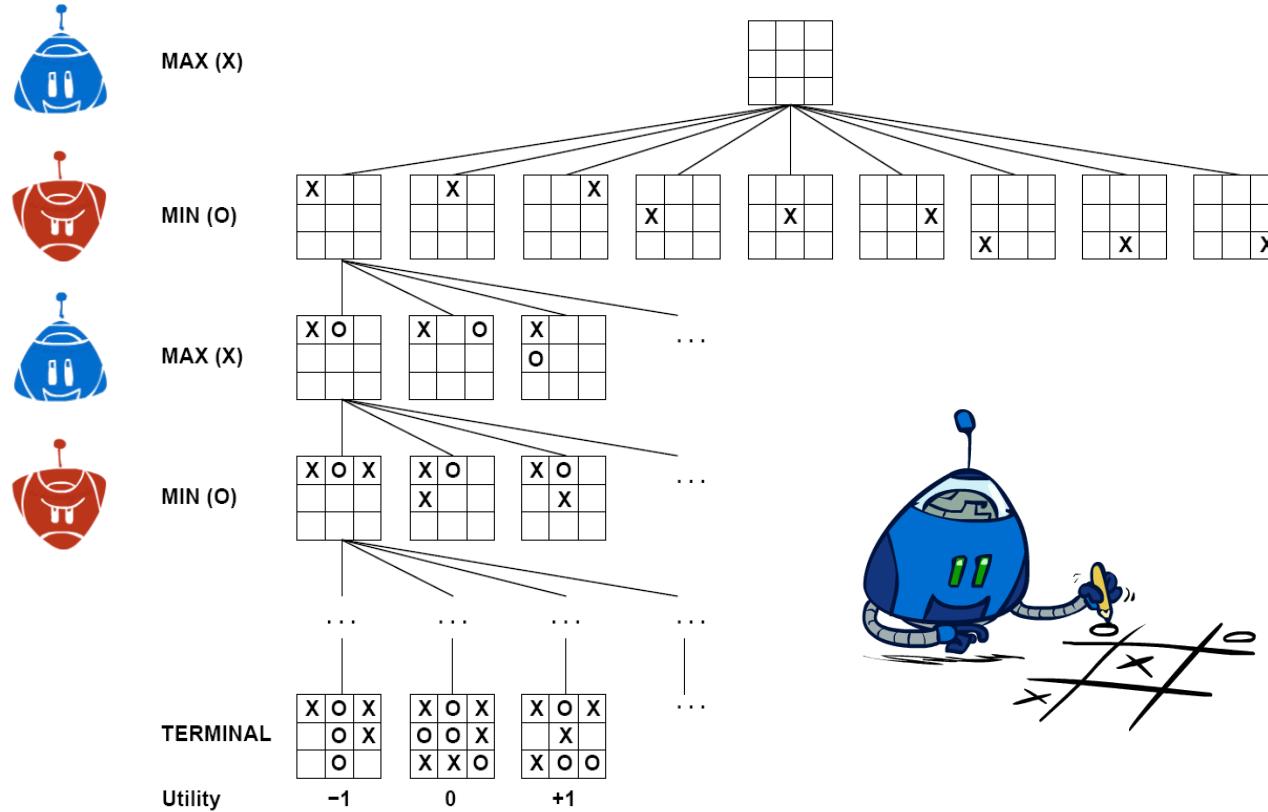
# Minimax Example



# Class work

- Halving game
  - Start with a number N.
  - Players take turns. Actions
    - Decrement number, or
    - Replace with  $\text{floor}(N/2)$  (e.g.,  $\text{floor}(5/2) = 2$ )
  - The player that makes it 0 wins
  - A win is worth +100 points, loss is -100
- Draw the game tree starting with N=5
- Calculate the minimax values at all nodes
  - In this game, don't be confused by the state of a node and its utility/value. The state just happens to also be a number in this game
- What is the optimal action?

# Tic-Tac-Toe Game Tree



# Minimax Implementation

```
def value(state):
```

    if the state is a terminal state: return the state's utility  
    if the next agent is MAX: return max-value(state)  
    if the next agent is MIN: return min-value(state)

```
def max-value(state):
```

    initialize  $v = -\infty$

    for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

    return  $v$

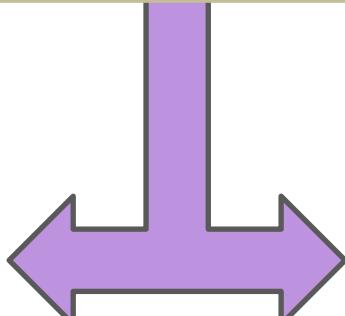
```
def min-value(state):
```

    initialize  $v = +\infty$

    for each successor of state:

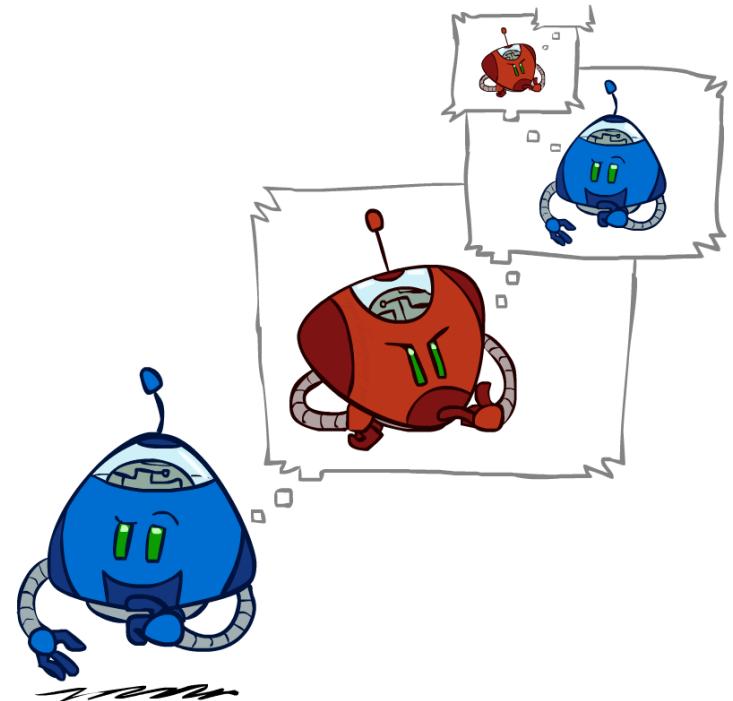
$v = \min(v, \text{value}(\text{successor}))$

    return  $v$

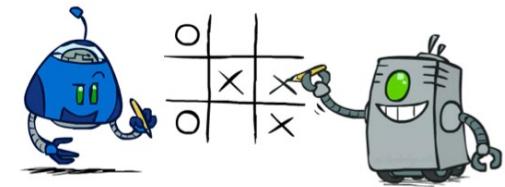
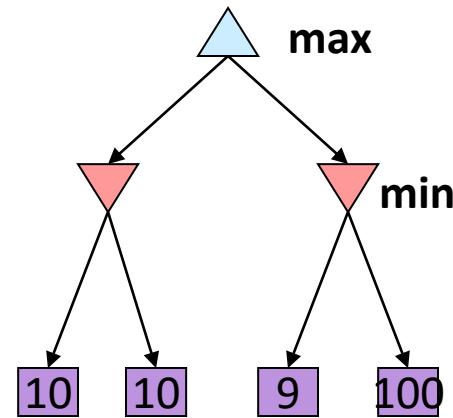
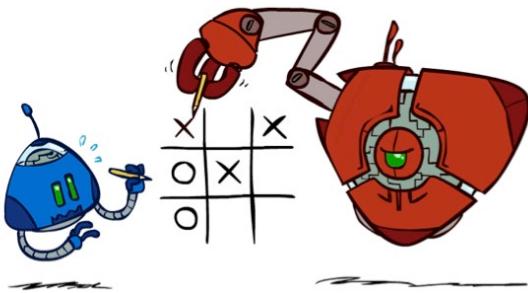


# Minimax Efficiency

- How efficient is minimax?
  - Just like (exhaustive) DFS
  - Time:  $O(b^m)$
  - Space:  $O(bm)$
- Example: For chess,  $b \approx 35$ ,  $m \approx 100$ 
  - Exact solution is completely infeasible
  - But, do we need to explore the whole tree?



# Minimax Properties



Optimal against a perfect player.

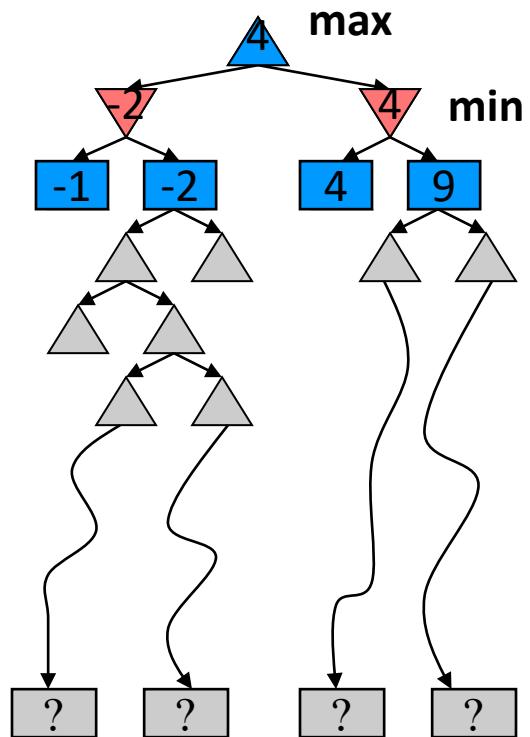
If player is known to be imperfect, other strategies could be better

# Resource Limits



# Speeding up minimax

- Problem: In realistic games, cannot search to leaves!
- Solution: Depth-limited search
  - Instead, search only to a limited depth in the tree
  - Replace terminal utilities with an **evaluation function** for non-terminal positions
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm



# Video: Limited Depth (2)

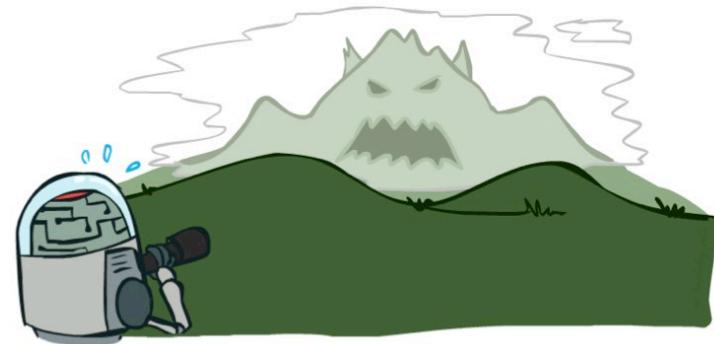


# Video: Limited Depth (10)

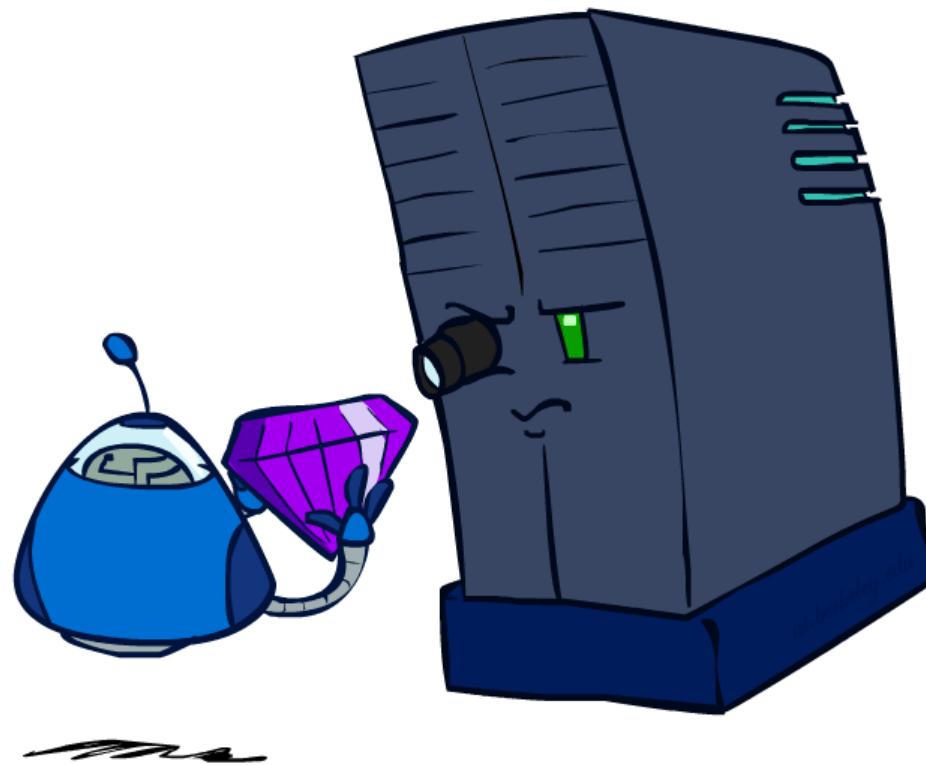


# Depth Matters

- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation

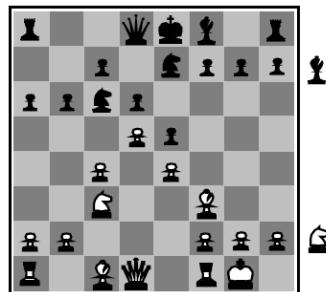


# Evaluation Functions



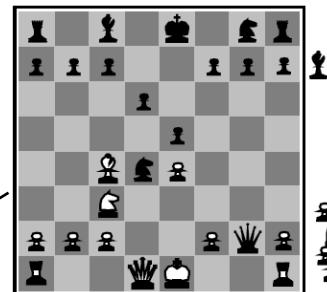
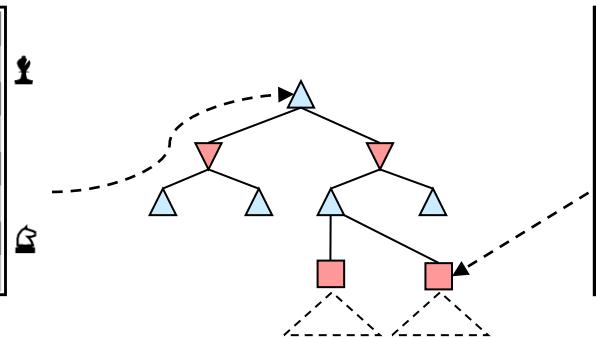
# Evaluation Functions

- Definition: An evaluation function is a (possibly very weak) **estimate** of the minimax value of an **internal** node
- Evaluation functions score non-terminals in depth-limited search



Black to move

White slightly better



White to move

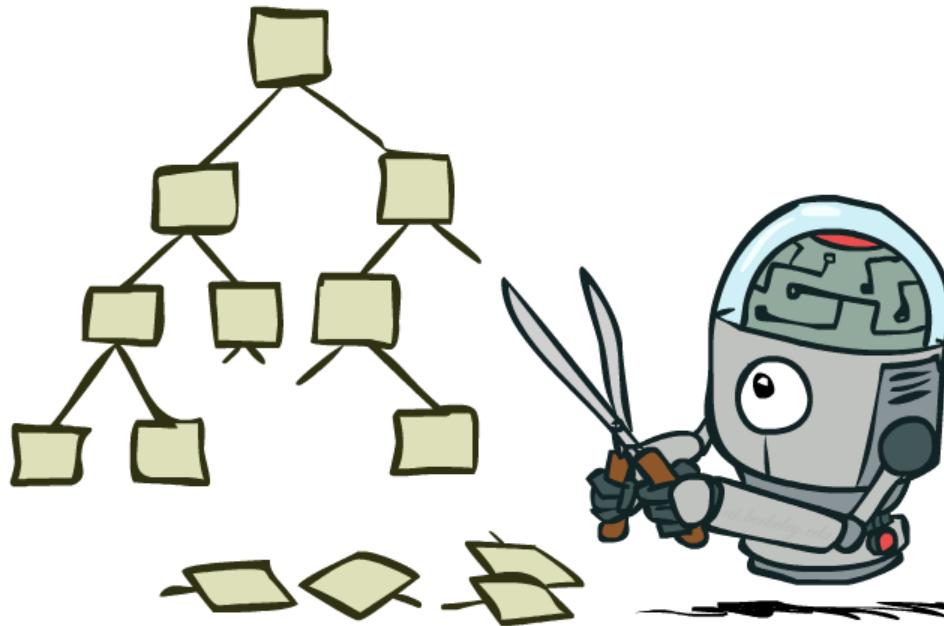
Black winning

- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

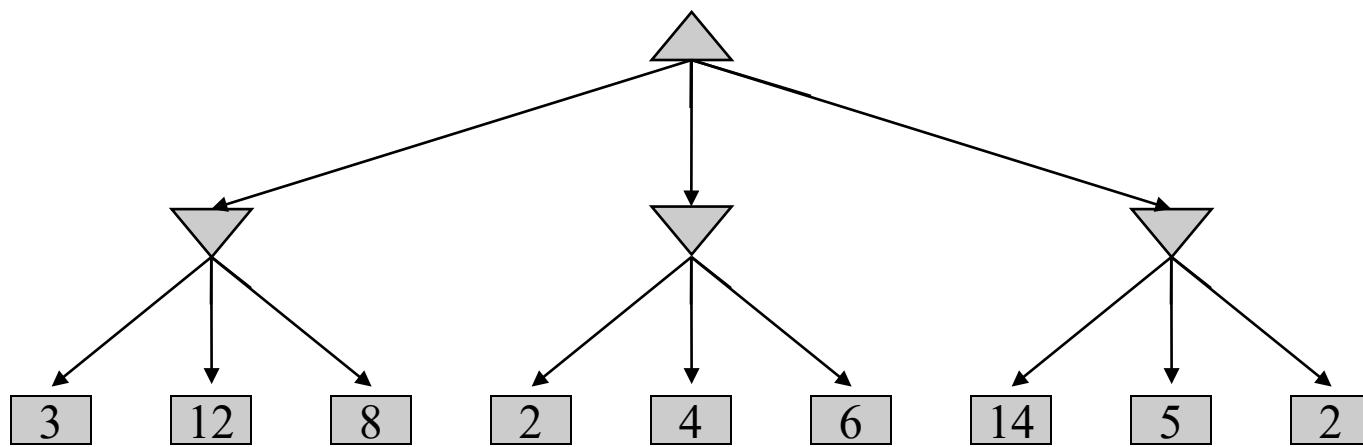
$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.  $f_1(s) = 9*(QW - QB) + 5*(RW - RB) + 3*(KW - KB) + 3*(BW - BB) + 1*(PW - PB)$

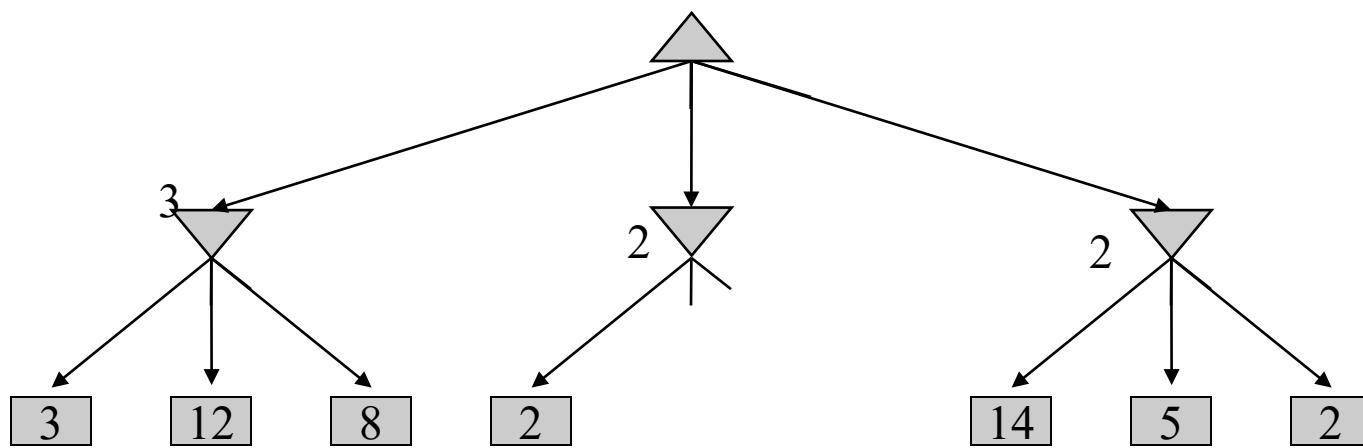
# Game Tree Pruning



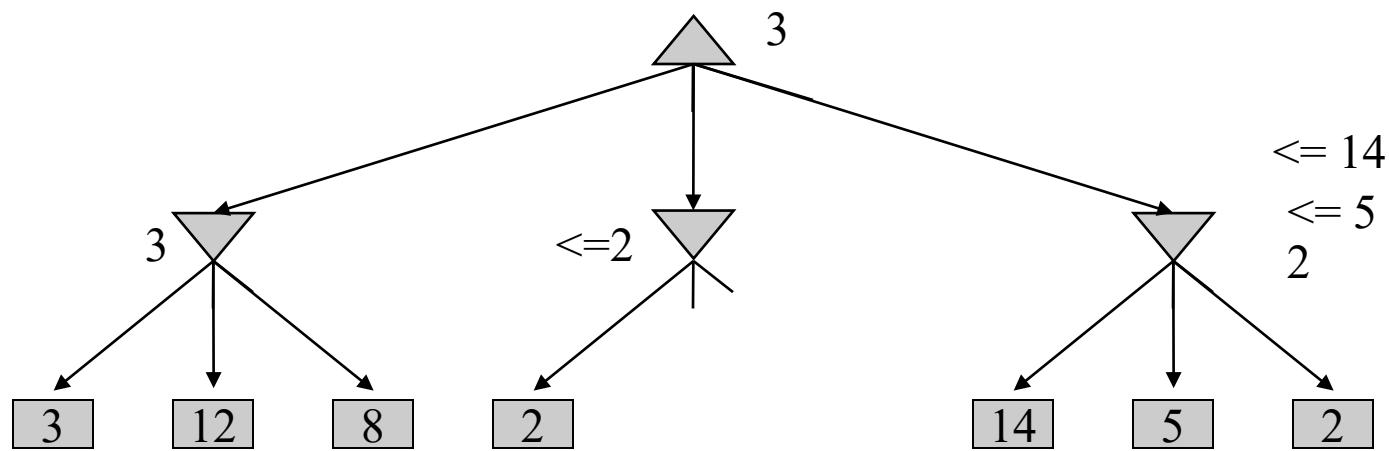
# Minimax Example



# Minimax Pruning

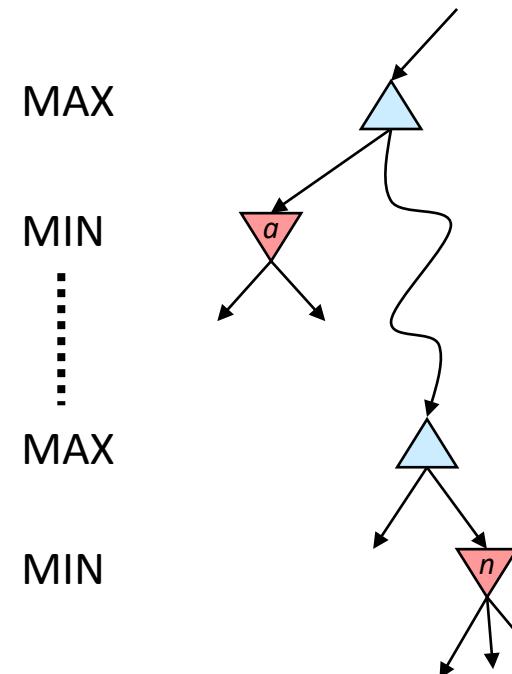


# Minimax Pruning



# Alpha-Beta Pruning

- General configuration (MIN version)
  - We're computing the MIN-VALUE at some node  $n$
  - We're looping over  $n$ 's children
  - $n$ 's estimate of the childrens' min is dropping
  - Who cares about  $n$ 's value? MAX
  - Let  $a$  be the best value that MAX can get at any choice point along the current path from the root
  - If  $n$  becomes worse than  $a$ , MAX will avoid it, so we can stop considering  $n$ 's other children (it's already bad enough that it won't be played)
- MAX version is symmetric



# Alpha-Beta Implementation

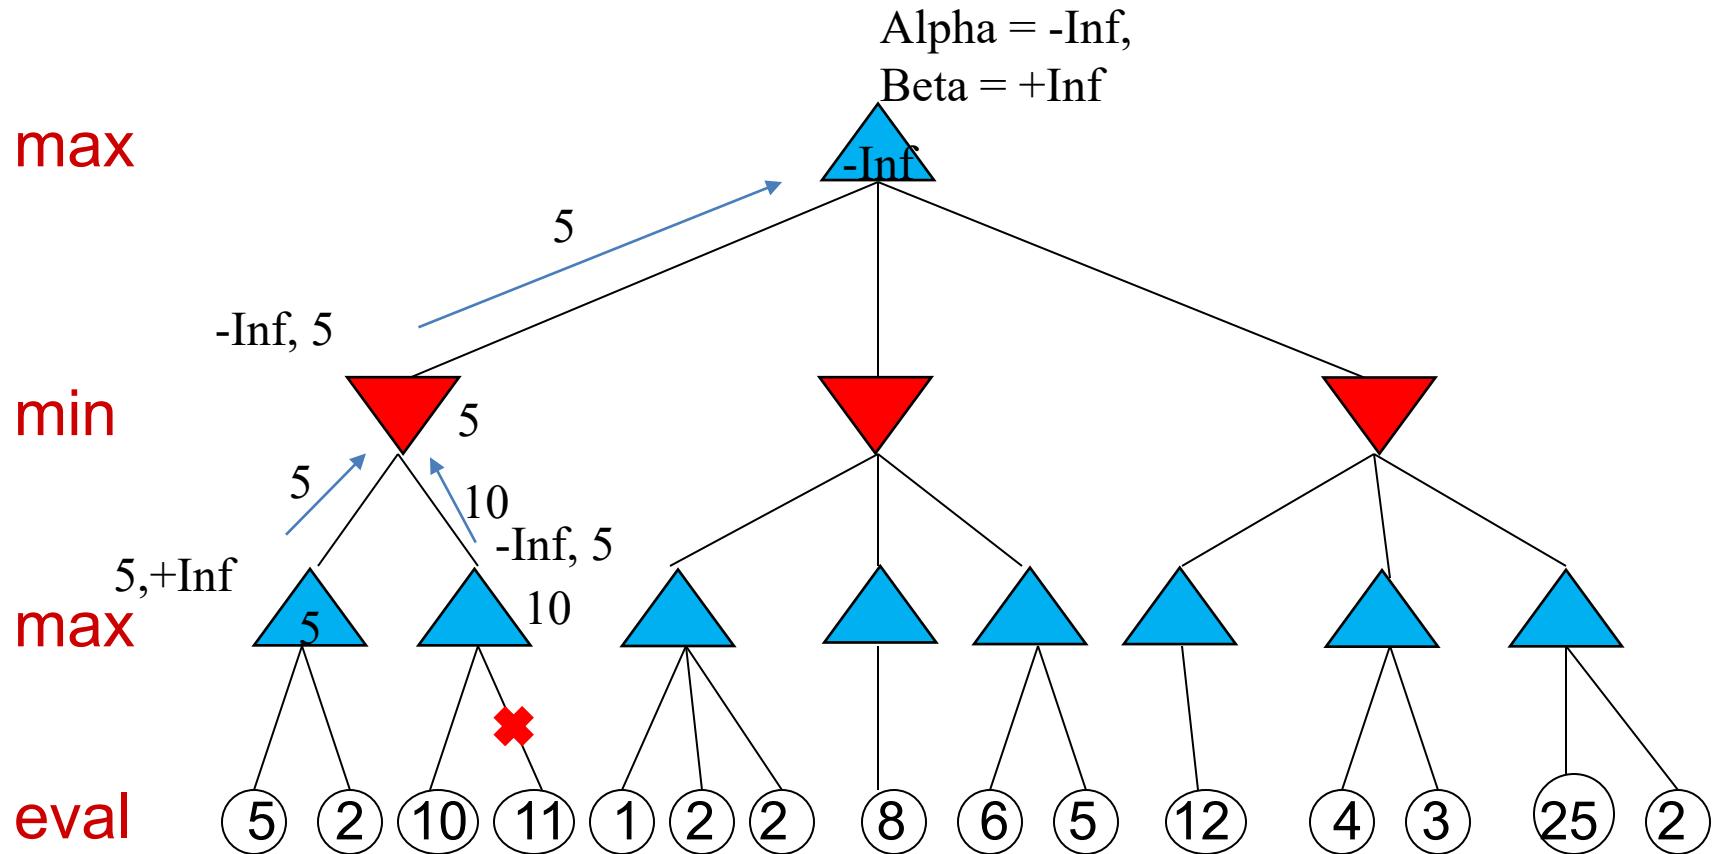
```
def value(state, α= -∞, β= +∞):  
    if the state is a terminal state: return the state's utility  
    if the next agent is MAX: return max-value(state)  
    if the next agent is MIN: return min-value(state)
```

```
def max-value(state, α, β):  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, value(successor, α, β))  
        if v ≥ β return v  
        α = max(α, v)  
    return v
```

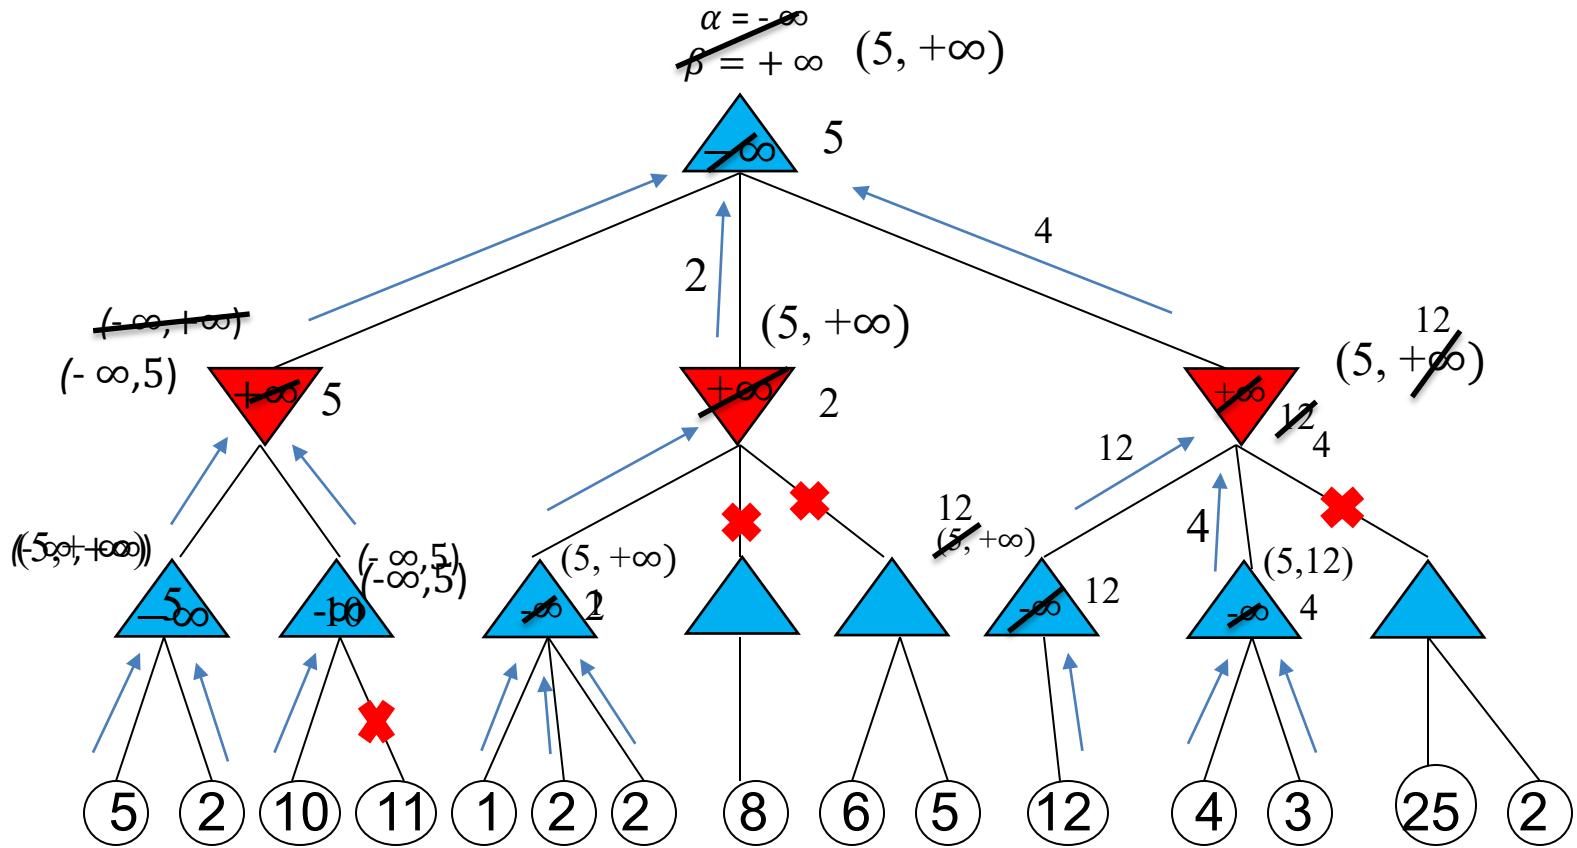
```
def min-value(state , α, β):  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, value(successor, α, β))  
        if v ≤ α return v  
        β = min(β, v)  
    return v
```

α: MAX's best option on path to root  
β: MIN's best option on path to root

# Alpha-Beta Pruning Practice

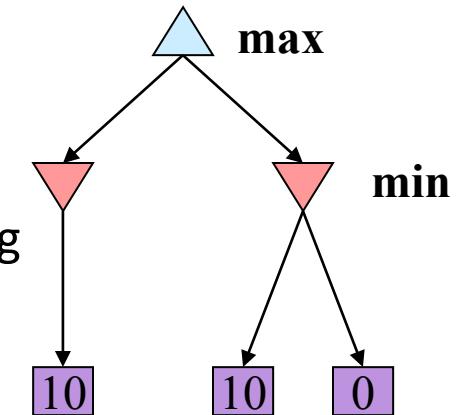


# Alpha-Beta Pruning Practice



# Alpha-Beta Pruning Properties

- This pruning has **no effect** on minimax value computed for the root!
  - No loss in accuracy
- Values of intermediate nodes might be wrong
  - Important: children of the root may have the wrong value
  - So the most naïve version won't let you do action selection
- Good **node ordering** improves effectiveness of pruning
  - Ideally, expand “most promising” children (moves) first
- With “perfect ordering”:
  - Time complexity drops to  $O(b^{m/2})$
  - Doubles solvable depth!
  - Full search of, e.g. chess, is still hopeless...



# Other adversarial search algorithms

- Monte Carlo tree search (MTCS)
  - Developed in 2006
  - MCTS was combined with neural networks in 2016 for computer Go
  - Instead of trying to search every path, randomly pick one and go down that path
    - And repeat many times
- Expectimax
  - A variant of minmax for games with randomness
  - E.g., a roll of a dice
  - MAX player now maximizes the *expected value* of the opponent's moves

# References

- George Luger, Artificial Intelligence: Structures and Strategies for Complex Problem Solving, 6<sup>th</sup> edition, Addison Wesley, 2009. **Chapter 4.4.**
- Russel and Norvig, Artificial Intelligence: A Modern Approach, 3<sup>rd</sup> edition, Prentice Hall, 2010. **Chapter 5.1-5.3.**

# Acknowledgement

- [http://ai.berkeley.edu/lecture slides.html](http://ai.berkeley.edu/lecture_slides.html)
- <https://www.youtube.com/watch?v=cwbjLlahbv8>